

```
!pip install opencv-python scikit-image n2v csbdeep
```

```
Requirement already satisfied: opencv-python in /usr/local/lib/python3.11/dist-packages (4.11.0.86)
Requirement already satisfied: scikit-image in /usr/local/lib/python3.11/dist-packages (0.25.2)
Collecting n2v
  Downloading n2v-0.3.3-py2.py3-none-any.whl.metadata (8.6 kB)
Collecting csbdeep
  Downloading csbdeep-0.8.1-py2.py3-none-any.whl.metadata (2.4 kB)
Requirement already satisfied: numpy>=1.21.2 in /usr/local/lib/python3.11/dist-packages (from opencv-python) (2.0.2)
Requirement already satisfied: scipy>=1.11.4 in /usr/local/lib/python3.11/dist-packages (from scikit-image) (1.15.2)
Requirement already satisfied: networkx>=3.0 in /usr/local/lib/python3.11/dist-packages (from scikit-image) (3.4.2)
Requirement already satisfied: pillow>=10.1 in /usr/local/lib/python3.11/dist-packages (from scikit-image) (11.2.1)
Requirement already satisfied: imageio!=2.35.0,>=2.33 in /usr/local/lib/python3.11/dist-packages (from scikit-image) (2.37.0)
Requirement already satisfied: tifffile>=2022.8.12 in /usr/local/lib/python3.11/dist-packages (from scikit-image) (2025.3.30)
Requirement already satisfied: packaging>=21 in /usr/local/lib/python3.11/dist-packages (from scikit-image) (24.2)
Requirement already satisfied: lazy-loader>=0.4 in /usr/local/lib/python3.11/dist-packages (from scikit-image) (0.4)
Collecting imagecodecs>=2020.2.18 (from n2v)
  Downloading imagecodecs-2025.3.30-cp311-cp311-manylinux_2_17_x86_64.manylinux2014_x86_64.whl.metadata (20 kB)
Collecting csbdeep
  Downloading csbdeep-0.7.4-py2.py3-none-any.whl.metadata (2.5 kB)
Collecting ruamel.yaml>=0.16.10 (from n2v)
  Downloading ruamel.yaml-0.18.10-py3-none-any.whl.metadata (23 kB)
Collecting bioimageio.core (from n2v)
  Downloading bioimageio_core-0.8.0-py3-none-any.whl.metadata (23 kB)
Requirement already satisfied: matplotlib in /usr/local/lib/python3.11/dist-packages (from csbdeep) (3.10.0)
Requirement already satisfied: six in /usr/local/lib/python3.11/dist-packages (from csbdeep) (1.17.0)
Requirement already satisfied: tqdm in /usr/local/lib/python3.11/dist-packages (from csbdeep) (4.67.1)
Collecting ruamel.yaml.clib>=0.2.7 (from ruamel.yaml>=0.16.10->n2v)
  Downloading ruamel.yaml.clib-0.2.12-cp311-cp311-manylinux_2_17_x86_64.manylinux2014_x86_64.whl.metadata (2.7 kB)
Collecting bioimageio.spec==0.5.4.1 (from bioimageio.core->n2v)
  Downloading bioimageio.spec-0.5.4.1-py3-none-any.whl.metadata (11 kB)
Requirement already satisfied: h5py in /usr/local/lib/python3.11/dist-packages (from bioimageio.core->n2v) (3.13.0)
Collecting loguru (from bioimageio.core->n2v)
  Downloading loguru-0.7.3-py3-none-any.whl.metadata (22 kB)
Collecting pydantic-settings<3,>=2.5 (from bioimageio.core->n2v)
  Downloading pydantic_settings-2.9.1-py3-none-any.whl.metadata (3.8 kB)
Requirement already satisfied: pydantic<3,>=2.7.0 in /usr/local/lib/python3.11/dist-packages (from bioimageio.core->n2v) (2.11.3)
Requirement already satisfied: requests in /usr/local/lib/python3.11/dist-packages (from bioimageio.core->n2v) (2.32.3)
Collecting ruyaml (from bioimageio.core->n2v)
  Downloading ruyaml-0.91.0-py3-none-any.whl.metadata (12 kB)
Requirement already satisfied: typing-extensions in /usr/local/lib/python3.11/dist-packages (from bioimageio.core->n2v) (4.13.2)
Collecting xarray<2025.3.0,>=2023.01 (from bioimageio.core->n2v)
  Downloading xarray-2025.1.2-py3-none-any.whl.metadata (11 kB)
Requirement already satisfied: annotated-types<1,>=0.5.0 in /usr/local/lib/python3.11/dist-packages (from bioimageio.spec==0.5.4.1->bioimageio.core->n2v) (0.7.0)
Collecting email-validator (from bioimageio.spec==0.5.4.1->bioimageio.core->n2v)
  Downloading email_validator-2.2.0-py3-none-any.whl.metadata (25 kB)
Requirement already satisfied: markdown in /usr/local/lib/python3.11/dist-packages (from bioimageio.spec==0.5.4.1->bioimageio.core->n2v) (3.7)
Requirement already satisfied: pooch<2,>=1.5 in /usr/local/lib/python3.11/dist-packages (from bioimageio.spec==0.5.4.1->bioimageio.core->n2v) (1.8.2)
Collecting pydantic<3,>=2.7.0 (from bioimageio.core->n2v)
  Downloading pydantic-2.9.2-py3-none-any.whl.metadata (149 kB)
149.4/149.4 kB 5.3 MB/s eta 0:00:00
Requirement already satisfied: python-dateutil in /usr/local/lib/python3.11/dist-packages (from bioimageio.spec==0.5.4.1->bioimageio.core->n2v) (2.9.0)
Requirement already satisfied: rich in /usr/local/lib/python3.11/dist-packages (from bioimageio.spec==0.5.4.1->bioimageio.core->n2v) (13.9.0)
Requirement already satisfied: zipp in /usr/local/lib/python3.11/dist-packages (from bioimageio.spec==0.5.4.1->bioimageio.core->n2v) (3.19.2)
Requirement already satisfied: contourpy>=1.0.1 in /usr/local/lib/python3.11/dist-packages (from matplotlib->csbdeep) (1.3.2)
Requirement already satisfied: cycler>=0.10 in /usr/local/lib/python3.11/dist-packages (from matplotlib->csbdeep) (0.12.1)
Requirement already satisfied: fonttools>=4.22.0 in /usr/local/lib/python3.11/dist-packages (from matplotlib->csbdeep) (4.57.0)
Requirement already satisfied: kiwisolver>=1.3.1 in /usr/local/lib/python3.11/dist-packages (from matplotlib->csbdeep) (1.4.8)
Requirement already satisfied: pyparsing>=2.3.1 in /usr/local/lib/python3.11/dist-packages (from matplotlib->csbdeep) (3.2.3)
```

```
!pip install PyWavelets # Install the missing PyWavelets package
```

```
Collecting PyWavelets
  Downloading pywavelets-1.8.0-cp311-cp311-manylinux_2_17_x86_64.manylinux2014_x86_64.whl.metadata (9.0 kB)
Requirement already satisfied: numpy<3,>=1.23 in /usr/local/lib/python3.11/dist-packages (from PyWavelets) (2.0.2)
Downloading pywavelets-1.8.0-cp311-cp311-manylinux_2_17_x86_64.manylinux2014_x86_64.whl (4.5 MB)
4.5/4.5 MB 37.6 MB/s eta 0:00:00
Installing collected packages: PyWavelets
Successfully installed PyWavelets-1.8.0
```

```
# Install required packages
```

```
!pip install scikit-image opencv-python-headless tensorflow matplotlib --quiet
```

```
import os
import cv2
import numpy as np
import matplotlib.pyplot as plt
from skimage import img_as_float
from skimage.restoration import denoise_wavelet
from skimage.metrics import peak_signal_noise_ratio, structural_similarity, mean_squared_error
from tensorflow.keras.models import Model
from tensorflow.keras.layers import Input, Conv2D, MaxPooling2D, UpSampling2D
from tensorflow.keras.optimizers import Adam
```



```

from tensorflow.keras.preprocessing.image import img_to_array, array_to_img

# Create folders
os.makedirs("original_images", exist_ok=True)
os.makedirs("denoised_results", exist_ok=True)

# Utility: Add synthetic noise to image
def add_noise(img, noise_type="gaussian"):
    row, col, ch = img.shape
    if noise_type == "gaussian":
        mean = 0
        sigma = 25
        gauss = np.random.normal(mean, sigma, (row, col, ch)).reshape(row, col, ch)
        noisy = img + gauss
        return np.clip(noisy, 0, 255).astype(np.uint8)
    return img

# Utility: Denoising Autoencoder
def build_denoising_autoencoder(input_shape):
    input_img = Input(shape=input_shape)
    x = Conv2D(32, (3, 3), activation='relu', padding='same')(input_img)
    x = MaxPooling2D((2, 2), padding='same')(x)
    x = Conv2D(16, (3, 3), activation='relu', padding='same')(x)
    encoded = MaxPooling2D((2, 2), padding='same')(x)

    x = Conv2D(16, (3, 3), activation='relu', padding='same')(encoded)
    x = UpSampling2D((2, 2))(x)
    x = Conv2D(32, (3, 3), activation='relu', padding='same')(x)
    x = UpSampling2D((2, 2))(x)
    decoded = Conv2D(3, (3, 3), activation='sigmoid', padding='same')(x)

    autoencoder = Model(input_img, decoded)
    autoencoder.compile(optimizer=Adam(), loss='mean_squared_error')
    return autoencoder

# Main processing function
def denoise_and_compare(image_path):
    img = cv2.imread(image_path)
    img_rgb = cv2.cvtColor(img, cv2.COLOR_BGR2RGB)
    img_resized = cv2.resize(img_rgb, (128, 128)) # Resize for DAE
    noisy_img = add_noise(img_resized).astype(np.uint8)

    # Convert to float
    img_float = img_as_float(noisy_img)

    # 1. Median Filter
    median = cv2.medianBlur(noisy_img, 3)

    # 2. Wavelet Denoising
    wavelet = denoise_wavelet(img_float, channel_axis=-1, rescale_sigma=True)
    wavelet_uint8 = (np.clip(wavelet, 0, 1) * 255).astype(np.uint8)

    # 3. Denoising Autoencoder
    x_train = np.array([noisy_img]) / 255.0
    y_train = np.array([img_resized]) / 255.0
    dae = build_denoising_autoencoder(input_shape=(128, 128, 3))
    dae.fit(x_train, y_train, epochs=100, verbose=0)
    dae_output = dae.predict(x_train)[0]
    dae_output_uint8 = (dae_output * 255).astype(np.uint8)

    # Save outputs
    cv2.imwrite("denoised_results/original.jpg", cv2.cvtColor(img_resized, cv2.COLOR_RGB2BGR))
    cv2.imwrite("denoised_results/noisy.jpg", cv2.cvtColor(noisy_img, cv2.COLOR_RGB2BGR))
    cv2.imwrite("denoised_results/median.jpg", cv2.cvtColor(median, cv2.COLOR_RGB2BGR))
    cv2.imwrite("denoised_results/wavelet.jpg", cv2.cvtColor(wavelet_uint8, cv2.COLOR_RGB2BGR))
    cv2.imwrite("denoised_results/dae.jpg", cv2.cvtColor(dae_output_uint8, cv2.COLOR_RGB2BGR))

    # Metric calculations
    def metrics(original, filtered):
        return {
            "PSNR": peak_signal_noise_ratio(original, filtered),
            "SSIM": structural_similarity(original, filtered, channel_axis=2),
            "MSE": mean_squared_error(original, filtered)
        }

    print("=== Image Denoising Metrics ===")
    print("Median Filter:", metrics(img_resized, median))
    print("Wavelet Filter:", metrics(img_resized, wavelet_uint8))
    print("Denoising Autoencoder:", metrics(img_resized, dae_output_uint8))

# Upload and process one image
from google.colab import files

```



```

uploaded = files.upload()
import shutil
for filename in uploaded:
    shutil.move(filename, f"original_images/{filename}")

```

```

# Denoise the uploaded image
img_path = os.listdir("original_images")[0]
denoise_and_compare(f"original_images/{img_path}")

```

Choose Files No file chosen Upload widget is only available when the cell has been executed in the current browser session. Please rerun this cell to enable.

Saving i3.jpg to i3 (1).jpg
 1/1 — 0s 213ms/step

=== Image Denoising Metrics ===
 Median Filter: {'PSNR': np.float64(25.133760693201673), 'SSIM': np.float64(0.717523015893685), 'MSE': np.float64(199.39042154947916)
 Wavelet Filter: {'PSNR': np.float64(25.5884035478262), 'SSIM': np.float64(0.741181142110028), 'MSE': np.float64(182.0072282877604)}

```


# Visualization of results
def show_results():
    titles = ['Original', 'Noisy', 'Median Filter', 'Wavelet Denoising', 'Denoising Autoencoder']
    files = ['original.jpg', 'noisy.jpg', 'median.jpg', 'wavelet.jpg', 'dae.jpg']
    plt.figure(figsize=(15, 8))

    for i, (title, file) in enumerate(zip(titles, files)):
        img = cv2.imread(f'denoised_results/{file}')
        img = cv2.cvtColor(img, cv2.COLOR_BGR2RGB)
        plt.subplot(2, 3, i+1)
        plt.imshow(img)
        plt.title(title)
        plt.axis('off')

    plt.tight_layout()
    plt.show()

# Call this function after denoising
show_results()

```



```

# Load image
import cv2
import numpy as np
import matplotlib.pyplot as plt

img_path = "/content/i3.jpg" # update with your actual filename
img = cv2.imread(img_path)
img_rgb = cv2.cvtColor(img, cv2.COLOR_BGR2RGB)
img_resized = cv2.resize(img_rgb, (128, 128)) # Resize for autoencoder

```



```
# Add synthetic Gaussian noise
noisy_img = add_noise(img_resized).astype(np.uint8)

# Normalize and prepare input
input_img = np.expand_dims(noisy_img / 255.0, axis=0)

# Build and train autoencoder (simulating Noise2Void behavior)
autoencoder = build_denoising_autoencoder((128, 128, 3))
autoencoder.fit(input_img, input_img, epochs=100, verbose=0) # Train on noisy image only

# Predict denoised output
denoised = autoencoder.predict(input_img)[0]
denoised_uint8 = (denoised * 255).astype(np.uint8)

# Save outputs
cv2.imwrite("denoised_results/noise2void_input.jpg", cv2.cvtColor(noisy_img, cv2.COLOR_RGB2BGR))
cv2.imwrite("denoised_results/noise2void_output.jpg", cv2.cvtColor(denoised_uint8, cv2.COLOR_RGB2BGR))

# Show comparison
plt.figure(figsize=(10, 4))
plt.subplot(1, 3, 1)
plt.title("Original")
plt.imshow(img_resized)
plt.axis('off')


plt.subplot(1, 3, 2)
plt.title("Noisy")
plt.imshow(noisy_img)
plt.axis('off')

plt.subplot(1, 3, 3)
plt.title("Noise2Void Output")
plt.imshow(denoised_uint8)
plt.axis('off')

plt.tight_layout()
plt.show()
```

1/1 — 0s 209ms/step


Original



Noisy



Noise2Void Output



```
# Install OpenCV if not already installed
!pip install opencv-python-headless --quiet
```

```
import cv2
import os
import numpy as np
import matplotlib.pyplot as plt
from google.colab import files
```

```
# Step 1: Upload video
uploaded = files.upload()
video_file = list(uploaded.keys())[0]
```

```
# Create output directories
os.makedirs("frames", exist_ok=True)
os.makedirs("processed_videos", exist_ok=True)
os.makedirs("collage_frames", exist_ok=True)
```

```
# Step 2: Extract Frames from Video
cap = cv2.VideoCapture(video_file)
frame_count = 0
frames = []
```




```

while True:
    ret, frame = cap.read()
    if not ret:
        break
    frame_path = f"frames/frame_{frame_count:04d}.jpg"
    cv2.imwrite(frame_path, frame)
    frames.append(frame)
    frame_count += 1

cap.release()
print(f"Total frames extracted: {frame_count}")

```

 Choose Files No file chosen Upload widget is only available when the cell has been executed in the current browser session. Please rerun this cell to enable.
Saving 0021800110-154.mp4 to 0021800110-154 (1).mp4

```

# Define video writer settings
height, width, layers = frames[0].shape
fourcc = cv2.VideoWriter_fourcc(*'mp4v')

# Define output writers
out_thresh = cv2.VideoWriter("processed_videos/adaptive_threshold.mp4", fourcc, 20.0, (width, height), False)
out_blur = cv2.VideoWriter("processed_videos/gaussian_blur.mp4", fourcc, 20.0, (width, height), True)
out_edges = cv2.VideoWriter("processed_videos/canny_edges.mp4", fourcc, 20.0, (width, height), False)
out_not = cv2.VideoWriter("processed_videos/bitwise_not.mp4", fourcc, 20.0, (width, height), True)

for idx, frame in enumerate(frames):
    gray = cv2.cvtColor(frame, cv2.COLOR_BGR2GRAY)

    # Adaptive Thresholding
    thresh = cv2.adaptiveThreshold(gray, 255,
                                   cv2.ADAPTIVE_THRESH_GAUSSIAN_C,
                                   cv2.THRESH_BINARY, 11, 2)

    out_thresh.write(thresh)

    # Gaussian Blur
    blur = cv2.GaussianBlur(frame, (15, 15), 0)
    out_blur.write(blur)

    # Canny Edge Detection
    edges = cv2.Canny(gray, 100, 200)
    out_edges.write(edges)


    # Bitwise NOT
    bitwise_not = cv2.bitwise_not(frame)
    out_not.write(bitwise_not)

    # Save some frames for collage
    if idx % 20 == 0 and idx <= 100:
        cv2.imwrite(f"collage_frames/frame_{idx}.jpg", frame)

# Release video writers
out_thresh.release()
out_blur.release()
out_edges.release()
out_not.release()

print("Processed videos saved in 'processed_videos/' folder.")

```

 Processed videos saved in 'processed_videos/' folder.

```

import glob

collage_images = sorted(glob.glob("collage_frames/*.jpg"))[:9]
collage_frames = [cv2.cvtColor(cv2.imread(img), cv2.COLOR_BGR2RGB) for img in collage_images]

fig, axes = plt.subplots(3, 3, figsize=(12, 12))
for ax, img in zip(axes.flatten(), collage_frames):
    ax.imshow(img)
    ax.axis('off')




plt.suptitle("Collage of Sample Frames", fontsize=16)
plt.tight_layout()
plt.show()

```





Collage of Sample Frames



```
import kagglehub
import os

# Step 1: Download the dataset
path = kagglehub.dataset_download("pevogam/ucf101")

# Step 2: Explore the downloaded files and directories
print("Base download path:", path)
print("\nContents of the downloaded path:")
print(os.listdir(path))

# Step 3: Search for the UCF101 video classes
for root, dirs, files in os.walk(path):
    print(f"\nFound directory: {root}")
    for d in dirs:
        print("  └", d)
    break # only print the first level
```



Base download path: /kaggle/input/ucf101

Contents of the downloaded path:
['UCF101', 'UCF101TrainTestSplits-RecognitionTask']

Found directory: /kaggle/input/ucf101
└ UCF101
└ UCF101TrainTestSplits-RecognitionTask



```
import os

root_dir = "/kaggle/input/ucf101/UCF101"
for root, dirs, files in os.walk(root_dir):
    print("Directory:", root)
    for d in dirs:
        print(" Subfolder:", d)
    for f in files[:3]: # print only a few files for brevity
        print(" File:", f)
    print("-" * 40)
```

```
↗ Directory: /kaggle/input/ucf101/UCF101
  Subfolder: UCF-101
-----
Directory: /kaggle/input/ucf101/UCF101/UCF-101
  Subfolder: HorseRace
  Subfolder: StillRings
  Subfolder: ApplyLipstick
  Subfolder: HammerThrow
  Subfolder: VolleyballSpiking
  Subfolder: Biking
  Subfolder: PlayingCello
  Subfolder: BodyWeightSquats
  Subfolder: TaiChi
  Subfolder: Punch
  Subfolder: BreastStroke
  Subfolder: Billiards
  Subfolder: BoxingPunchingBag
  Subfolder: BasketballDunk
  Subfolder: PoleVault
  Subfolder: ThrowDiscus
  Subfolder: BaseballPitch
  Subfolder: Knitting
  Subfolder: SumoWrestling
  Subfolder: HorseRiding
  Subfolder: Mixing
  Subfolder: BrushingTeeth
  Subfolder: HighJump
  Subfolder: Skijet
  Subfolder: SkateBoarding
  Subfolder: MilitaryParade
  Subfolder: IceDancing
  Subfolder: CricketShot
  Subfolder: Fencing
  Subfolder: JugglingBalls
  Subfolder: Swing
  Subfolder: RockClimbingIndoor
  Subfolder: PlayingFlute
  Subfolder: SalsaSpin
  Subfolder: CricketBowling
  Subfolder: Typing
  Subfolder: ApplyEyeMakeup
  Subfolder: PlayingTabla
  Subfolder: BalanceBeam
  Subfolder: FloorGymnastics
  Subfolder: HeadMassage
  Subfolder: FrisbeeCatch
  Subfolder: Rowing
  Subfolder: Hammering
  Subfolder: CuttingInKitchen
  Subfolder: BenchPress
  Subfolder: PushUps
  Subfolder: Nunchucks
  Subfolder: Archery
  Subfolder: LongJump
  Subfolder: BlowingCandles
  Subfolder: WallPushups
  Subfolder: PlayingViolin
  Subfolder: PullUps
```

```
import os
import shutil
import random

# Define source and destination directories
SOURCE_DIR = '/kaggle/input/ucf101/UCF101/UCF-101'
DEST_DIR = '/kaggle/working/UCF101_subset'

# List of selected classes (can be updated as needed)
SELECTED_CLASSES = ['Basketball', 'Biking', 'PlayingGuitar', 'Typing', 'JumpRope']
VIDEOS_PER_CLASS = 10

# Create the destination directory if it doesn't exist
os.makedirs(DEST_DIR, exist_ok=True)

# Iterate over the selected classes and copy videos
```



```

for cls in SELECTED_CLASSES:
    class_path = os.path.join(SOURCE_DIR, cls)
    dest_class_path = os.path.join(DEST_DIR, cls)


    # Create class folder in destination
    os.makedirs(dest_class_path, exist_ok=True)

    # Select random 10 videos from the class
    selected = random.sample(os.listdir(class_path), VIDEOS_PER_CLASS)

    # Copy selected videos to the destination
    for video in selected:
        shutil.copy(os.path.join(class_path, video), dest_class_path)

print(f"Subset created at: {DEST_DIR}")

```

 Subset created at: /kaggle/working/UCF101_subset

```

import cv2
import os
import numpy as np

# Define parameters
FRAME_RATE = 5 # Extract every 5th frame
RESIZE_DIM = (112, 112) # Resize frames to 112x112
MAX_FRAMES = 16 # Number of frames per video

# Function to extract frames from video
def extract_frames(video_path, max_frames=MAX_FRAMES, frame_rate=FRAME_RATE, resize_dim=RESIZE_DIM):
    # Read the video
    cap = cv2.VideoCapture(video_path)

    frames = []
    frame_count = 0
    while True:
        ret, frame = cap.read()
        if not ret:
            break

        # Extract every 'frame_rate'-th frame
        if frame_count % frame_rate == 0:
            frame_resized = cv2.resize(frame, resize_dim)
            frames.append(frame_resized)

        frame_count += 1

    # Stop once we have extracted enough frames
    if len(frames) == max_frames:
        break

    cap.release()

    # If fewer than MAX_FRAMES are extracted, pad the sequence with the last frame
    while len(frames) < max_frames:
        frames.append(frames[-1])

    return np.array(frames)

# Process each class and video
video_frames = {}
for cls in SELECTED_CLASSES:
    class_path = os.path.join(DEST_DIR, cls)
    video_frames[cls] = []

    for video in os.listdir(class_path):
        video_path = os.path.join(class_path, video)
        frames = extract_frames(video_path)
        video_frames[cls].append(frames)

print("Frame extraction completed.")

```

 Frame extraction completed.

```

from sklearn.preprocessing import LabelEncoder

# Initialize label encoder
label_encoder = LabelEncoder()

# Fit the encoder on the selected classes

```




```
labels = label_encoder.fit_transform(SELECTED_CLASSES)
```

```
# Create a dictionary to map class names to labels
class_labels = dict(zip(SELECTED_CLASSES, labels))
print("Label encoding completed.")
```

↗ Label encoding completed.

```
from sklearn.model_selection import train_test_split
```

```
# Prepare data and labels
data = []
labels = []
```

```
# Add frames and their corresponding labels
for cls in SELECTED_CLASSES:
    for frames in video_frames[cls]:
        data.append(frames)
        labels.append(class_labels[cls])
```

```
# Convert to numpy arrays
data = np.array(data)
labels = np.array(labels)
```

```
# Split data into training and testing sets (80/20 split)
X_train, X_test, y_train, y_test = train_test_split(data, labels, test_size=0.2, random_state=42)
```

```
print(f"Training data shape: {X_train.shape}")
print(f"Testing data shape: {X_test.shape}")
```

↗ Training data shape: (40, 16, 112, 112, 3)
Testing data shape: (10, 16, 112, 112, 3)

```
import tensorflow as tf
from tensorflow.keras import layers, models
```

```
# Define 3D CNN Model
```

```
def create_3d_cnn_model(input_shape, num_classes):
    model = models.Sequential()
```

```
    # 3D convolution layers with padding='same'
    model.add(layers.Conv3D(32, kernel_size=(3, 3, 3), activation='relu', input_shape=input_shape, padding='same'))
    model.add(layers.MaxPooling3D(pool_size=(2, 2, 2)))
    model.add(layers.Conv3D(64, kernel_size=(3, 3, 3), activation='relu', padding='same'))
    model.add(layers.MaxPooling3D(pool_size=(2, 2, 2)))
    model.add(layers.Conv3D(128, kernel_size=(3, 3, 3), activation='relu', padding='same'))
    model.add(layers.MaxPooling3D(pool_size=(2, 2, 2)))
```

```
    # Flatten and fully connected layers
    model.add(layers.Flatten())
    model.add(layers.Dense(128, activation='relu'))
    model.add(layers.Dense(num_classes, activation='softmax'))
```

```
    return model
```

```
# Define input shape based on frame size and sequence length
input_shape = (MAX_FRAMES, RESIZE_DIM[0], RESIZE_DIM[1], 3) # (16, 112, 112, 3) for 16 frames of 112x112 images with 3 color channels
num_classes = len(SELECTED_CLASSES)
```

```
# Create the model
model = create_3d_cnn_model(input_shape, num_classes)
```

```
# Compile the model
model.compile(optimizer='adam', loss='sparse_categorical_crossentropy', metrics=['accuracy'])
```

```
# Model summary
model.summary()
```



```

/usr/local/lib/python3.11/dist-packages/keras/src/layers/convolutional/base_conv.py:107: UserWarning: Do not pass an `input_shape` to
super().__init__(activity_regularizer=activity_regularizer, **kwargs)
Model: "sequential_1"

```

Layer (type)	Output Shape	Param #
conv3d_3 (Conv3D)	(None, 16, 112, 112, 32)	2,624
max_pooling3d_3 (MaxPooling3D)	(None, 8, 56, 56, 32)	0
conv3d_4 (Conv3D)	(None, 8, 56, 56, 64)	55,360
max_pooling3d_4 (MaxPooling3D)	(None, 4, 28, 28, 64)	0
conv3d_5 (Conv3D)	(None, 4, 28, 28, 128)	221,312
max_pooling3d_5 (MaxPooling3D)	(None, 2, 14, 14, 128)	0
flatten_1 (Flatten)	(None, 50176)	0
dense_2 (Dense)	(None, 128)	6,422,656
dense_3 (Dense)	(None, 5)	645

Total params: 6,702,597 (25.57 MB)
Trainable params: 6,702,597 (25.57 MB)

```

# Train the model
history = model.fit(X_train, y_train, epochs=10, batch_size=4, validation_data=(X_test, y_test))

# Save the model
model.save("/kaggle/working/ucf101_3dcnn_model.h5")

print("Training completed and model saved.")

```

```

Epoch 1/10
10/10 — 53s 5s/step - accuracy: 0.1854 - loss: 398.4858 - val_accuracy: 0.4000 - val_loss: 1.5153
Epoch 2/10
10/10 — 81s 5s/step - accuracy: 0.2454 - loss: 1.5007 - val_accuracy: 0.2000 - val_loss: 2.3985
Epoch 3/10
10/10 — 52s 5s/step - accuracy: 0.5691 - loss: 1.3393 - val_accuracy: 0.1000 - val_loss: 1.5773
Epoch 4/10
10/10 — 52s 5s/step - accuracy: 0.6785 - loss: 1.1525 - val_accuracy: 0.3000 - val_loss: 2.2342
Epoch 5/10
10/10 — 80s 5s/step - accuracy: 0.6778 - loss: 1.4211 - val_accuracy: 0.3000 - val_loss: 2.1582
Epoch 6/10
10/10 — 83s 5s/step - accuracy: 0.5804 - loss: 0.9205 - val_accuracy: 0.1000 - val_loss: 1.7314
Epoch 7/10
10/10 — 83s 5s/step - accuracy: 0.8245 - loss: 0.6209 - val_accuracy: 0.3000 - val_loss: 1.7927
Epoch 8/10
10/10 — 82s 5s/step - accuracy: 0.7939 - loss: 0.5056 - val_accuracy: 0.6000 - val_loss: 2.0012
Epoch 9/10
10/10 — 79s 5s/step - accuracy: 0.8435 - loss: 0.4716 - val_accuracy: 0.5000 - val_loss: 1.4895
Epoch 10/10
10/10 — 82s 5s/step - accuracy: 0.7711 - loss: 0.5867 - val_accuracy: 0.5000 - val_loss: 1.6238
WARNING:absl:You are saving your model as an HDF5 file via `model.save()` or `keras.saving.save_model(model)`. This file format is
Training completed and model saved.

```

```

from sklearn.metrics import confusion_matrix, classification_report
import seaborn as sns
import matplotlib.pyplot as plt

# Evaluate the model on the test set
y_pred = model.predict(X_test)
y_pred_classes = np.argmax(y_pred, axis=1)

# Print classification report
print("Classification Report:")
print(classification_report(y_test, y_pred_classes, target_names=SELECTED_CLASSES))

# Compute confusion matrix
cm = confusion_matrix(y_test, y_pred_classes)

# Plot confusion matrix
plt.figure(figsize=(8, 6))
sns.heatmap(cm, annot=True, fmt='d', cmap='Blues', xticklabels=SELECTED_CLASSES, yticklabels=SELECTED_CLASSES)
plt.xlabel('Predicted')
plt.ylabel('True')
plt.title('Confusion Matrix')

```



plt.show()

1/1 3s 3s/step

Classification Report:

	precision	recall	f1-score	support
Basketball	0.00	0.00	0.00	0
Biking	0.67	0.67	0.67	3
PlayingGuitar	0.50	0.50	0.50	2
Typing	0.50	0.50	0.50	2
JumpRope	1.00	0.33	0.50	3
accuracy			0.50	10
macro avg	0.53	0.40	0.43	10
weighted avg	0.70	0.50	0.55	10

```

/usr/local/lib/python3.11/dist-packages/sklearn/metrics/_classification.py:1565: UndefinedMetricWarning: Recall is ill-defined and t
_warn_prf(average, modifier, f"{metric.capitalize()} is", len(result))
/usr/local/lib/python3.11/dist-packages/sklearn/metrics/_classification.py:1565: UndefinedMetricWarning: Recall is ill-defined and t
_warn_prf(average, modifier, f"{metric.capitalize()} is", len(result))
/usr/local/lib/python3.11/dist-packages/sklearn/metrics/_classification.py:1565: UndefinedMetricWarning: Recall is ill-defined and t
_warn_prf(average, modifier, f"{metric.capitalize()} is", len(result))

```

